

RNA-Seq Workflow Template

Author: *Daniela Cassol (danielac@ucr.edu) and Thomas Girke (thomas.girke@ucr.edu)*

Last update: 29 October, 2020

Package

systemPipeR 1.25.1

Contents

1	Introduction	3
2	Samples and environment settings.	3
2.1	Environment settings and input data.	3
2.2	Required packages and resources	3
2.3	Experiment definition provided by <code>targets</code> file	3
3	Read preprocessing	4
3.1	Read quality filtering and trimming	4
3.2	FASTQ quality report	5
4	Alignments	5
4.1	Read mapping with <code>HISAT2</code>	5
4.2	Read and alignment stats.	8
4.3	Create symbolic links for viewing BAM files in IGV	8
5	Read quantification.	8
5.1	Read counting with <code>summarizeOverlaps</code> in parallel mode using multiple cores	8
5.2	Sample-wise correlation analysis	9
6	Analysis of DEGs.	10
6.1	Run <code>edgeR</code>	11
6.2	Plot DEG results	11
6.3	Venn diagrams of DEG sets.	12
7	GO term enrichment analysis.	13
7.1	Obtain gene-to-GO mappings.	13

RNA-Seq Workflow Template

- [7.2 Batch GO term enrichment analysis](#) 14
 - [7.3 Plot batch GO term results](#) 15
 - [8 Clustering and heat maps](#) 15
 - [9 Version Information](#). 16
 - [10 Funding](#) 18
 - [References](#) 18

1 Introduction

Users want to provide here background information about the design of their RNA-Seq project.

2 Samples and environment settings

2.1 Environment settings and input data

Typically, the user wants to record here the sources and versions of the reference genome sequence along with the corresponding annotations. In the provided sample data set all data inputs are stored in a `data` subdirectory and all results will be written to a separate `results` directory, while the `systemPipeRNAseq.Rmd` script and the `targets` file are expected to be located in the parent directory. The R session is expected to run from this parent directory.

`systemPipeRdata` package is a helper package to generate a fully populated `systemPipeR` workflow environment in the current working directory with a single command. All the instruction for generating the workflow are provide in the `systemPipeRdata` vignette [here](#).

The mini sample FASTQ files used by this report as well as the associated reference genome files can be loaded via the `systemPipeRdata` package. The chosen data set [SRP010938](#) contains 18 paired-end (PE) read sets from *Arabidopsis thaliana* (Howard et al. 2013). To minimize processing time during testing, each FASTQ file has been subsetting to 90,000-100,000 randomly sampled PE reads that map to the first 100,000 nucleotides of each chromosome of the *A. thaliana* genome. The corresponding reference genome sequence (FASTA) and its GFF annotation files have been truncated accordingly. This way the entire test sample data set is less than 200MB in storage space. A PE read set has been chosen for this test data set for flexibility, because it can be used for testing both types of analysis routines requiring either SE (single end) reads or PE reads.

2.2 Required packages and resources

The `systemPipeR` package needs to be loaded to perform the analysis steps shown in this report (H Backman and Girke 2016).

```
library(systemPipeR)
```

To apply workflows to custom data, the user needs to modify the `targets` file and if necessary update the corresponding parameter (`.cwl` and `.yml`) files. A collection of pre-generated `.cwl` and `.yml` files are provided in the `param/cwl` subdirectory of each workflow template. They are also viewable in the GitHub repository of `systemPipeRdata` ([see here](#)). For more information of the structure of the `targets` file, please consult the documentation [here](#). More details about the new parameter files from `systemPipeR` can be found [here](#).

2.3 Experiment definition provided by `targets` file

The `targets` file defines all FASTQ files and sample comparisons of the analysis workflow.

```
targetspath <- system.file("extdata", "targetsPE.txt", package = "systemPipeR")
targets <- read.delim(targetspath, comment.char = "#")[, 1:4]
targets
##                               FileName1                               FileName2
```

```
## 1 ./data/SRR446027_1.fastq.gz ./data/SRR446027_2.fastq.gz
## 2 ./data/SRR446028_1.fastq.gz ./data/SRR446028_2.fastq.gz
## 3 ./data/SRR446029_1.fastq.gz ./data/SRR446029_2.fastq.gz
## 4 ./data/SRR446030_1.fastq.gz ./data/SRR446030_2.fastq.gz
## 5 ./data/SRR446031_1.fastq.gz ./data/SRR446031_2.fastq.gz
## 6 ./data/SRR446032_1.fastq.gz ./data/SRR446032_2.fastq.gz
## 7 ./data/SRR446033_1.fastq.gz ./data/SRR446033_2.fastq.gz
## 8 ./data/SRR446034_1.fastq.gz ./data/SRR446034_2.fastq.gz
## 9 ./data/SRR446035_1.fastq.gz ./data/SRR446035_2.fastq.gz
## 10 ./data/SRR446036_1.fastq.gz ./data/SRR446036_2.fastq.gz
## 11 ./data/SRR446037_1.fastq.gz ./data/SRR446037_2.fastq.gz
## 12 ./data/SRR446038_1.fastq.gz ./data/SRR446038_2.fastq.gz
## 13 ./data/SRR446039_1.fastq.gz ./data/SRR446039_2.fastq.gz
## 14 ./data/SRR446040_1.fastq.gz ./data/SRR446040_2.fastq.gz
## 15 ./data/SRR446041_1.fastq.gz ./data/SRR446041_2.fastq.gz
## 16 ./data/SRR446042_1.fastq.gz ./data/SRR446042_2.fastq.gz
## 17 ./data/SRR446043_1.fastq.gz ./data/SRR446043_2.fastq.gz
## 18 ./data/SRR446044_1.fastq.gz ./data/SRR446044_2.fastq.gz
## SampleName Factor
## 1 M1A M1
## 2 M1B M1
## 3 A1A A1
## 4 A1B A1
## 5 V1A V1
## 6 V1B V1
## 7 M6A M6
## 8 M6B M6
## 9 A6A A6
## 10 A6B A6
## 11 V6A V6
## 12 V6B V6
## 13 M12A M12
## 14 M12B M12
## 15 A12A A12
## 16 A12B A12
## 17 V12A V12
## 18 V12B V12
```

3 Read preprocessing

3.1 Read quality filtering and trimming

The function `preprocessReads` allows to apply predefined or custom read preprocessing functions to all FASTQ files referenced in a `SYSargs2` container, such as quality filtering or adapter trimming routines. The paths to the resulting output FASTQ files are stored in the `output` slot of the `SYSargs2` object. The following example performs adapter trimming with the `trimLRPatterns` function from the `Biostrings` package. After the trimming step a new targets file is generated (here `targets_trim.txt`) containing the paths to the trimmed FASTQ files. The new targets file can be used for the next workflow step with an updated `SYSargs2` instance, e.g. running the NGS alignments using the trimmed FASTQ files.

RNA-Seq Workflow Template

Construct *SYSargs2* object from *cwl* and *yml* param and *targets* files.

```
dir_path <- system.file("extdata/cwl/preprocessReads/trim-pe",
  package = "systemPipeR")
trim <- loadWorkflow(targets = targetspath, wf_file = "trim-pe.cwl",
  input_file = "trim-pe.yml", dir_path = dir_path)
trim <- renderWF(trim, inputvars = c(FileName1 = "_FASTQ_PATH1_",
  FileName2 = "_FASTQ_PATH2_", SampleName = "_SampleName_"))
trim
output(trim)[1:2]

preprocessReads(args = trim, Fct = "trimLRPatterns(Rpattern='GCCCCGGGTAA',
  subject=fq)",
  batchsize = 1e+05, overwrite = TRUE, compress = TRUE)
writeTargetsout(x = trim, file = "targets_trim.txt", step = 1,
  new_col = c("FileName1", "FileName2"), new_col_output_index = c(1,
  2), overwrite = TRUE)
```

3.2 FASTQ quality report

The following *seeFastq* and *seeFastqPlot* functions generate and plot a series of useful quality statistics for a set of FASTQ files including per cycle quality box plots, base proportions, base-level quality trends, relative k-mer diversity, length and occurrence distribution of reads, number of reads above quality cutoffs and mean quality distribution. The results are written to a PDF file named *fastqReport.pdf*.

```
fqlist <- seeFastq(fastq = infile1(trim), batchsize = 10000,
  klength = 8)
pdf("./results/fastqReport.pdf", height = 18, width = 4 * length(fqlist))
seeFastqPlot(fqlist)
dev.off()
```




Figure 1: FASTQ quality report for 18 samples

4 Alignments

4.1 Read mapping with HISAT2

The following steps will demonstrate how to use the short read aligner *Hisat2* (Kim, Langmead, and Salzberg 2015) in both interactive job submissions and batch submissions to queuing systems of clusters using the *systemPipeR*'s new CWL command-line interface.

RNA-Seq Workflow Template

Build `Hisat2` index.

```
dir_path <- system.file("extdata/cwl/hisat2/hisat2-idx", package = "systemPipeR")
idx <- loadWorkflow(targets = NULL, wf_file = "hisat2-index.cwl",
  input_file = "hisat2-index.yml", dir_path = dir_path)
idx <- renderWF(idx)
idx
cmdlist(idx)

## Run
runCommandLine(idx, make_bam = FALSE)
```

The parameter settings of the aligner are defined in the `hisat2-mapping-se.cwl` and `hisat2-mapping-se.yml` files. The following shows how to construct the corresponding `SYSargs2` object, here `args`.

```
dir_path <- system.file("extdata/cwl/hisat2/hisat2-pe", package = "systemPipeR")
args <- loadWorkflow(targets = targetspath, wf_file = "hisat2-mapping-pe.cwl",
  input_file = "hisat2-mapping-pe.yml", dir_path = dir_path)
args <- renderWF(args, inputvars = c(FileName1 = "_FASTQ_PATH1_",
  FileName2 = "_FASTQ_PATH2_", SampleName = "_SampleName_"))
args
## Instance of 'SYSargs2':
##   Slot names/accessors:
##     targets: 18 (M1A...V12B), targetsheader: 4 (lines)
##     modules: 1
##     wf: 0, clt: 1, yamlinput: 8 (components)
##     input: 18, output: 18
##     cmdlist: 18
##   WF Steps:
##     1. hisat2-mapping-pe (rendered: TRUE)
cmdlist(args)[1:2]
## $M1A
## $M1A$`hisat2-mapping-pe`
## [1] "hisat2 -S ./results/M1A.sam -x ./data/tair10.fasta -k 1 --min-intronlen 30 --max-intronlen 3000"
##
## $M1B
## $M1B$`hisat2-mapping-pe`
## [1] "hisat2 -S ./results/M1B.sam -x ./data/tair10.fasta -k 1 --min-intronlen 30 --max-intronlen 3000"
output(args)[1:2]
## $M1A
## $M1A$`hisat2-mapping-pe`
## [1] "./results/M1A.sam"
##
## $M1B
## $M1B$`hisat2-mapping-pe`
## [1] "./results/M1B.sam"
```

4.1.1 Interactive job submissions in a single machine

To simplify the short read alignment execution for the user, the command-line can be run with the `runCommandLine` function. The execution will be on a single machine without submitting to a queuing system of a computer cluster. This way, the input FASTQ files will be processed sequentially. By default `runCommandLine` auto detects SAM file outputs and converts them to sorted and indexed BAM files, using internally the `Rsamtools` package (Morgan et al. 2019). Besides, `runCommandLine` allows the user to create a dedicated results folder for each workflow and a sub-folder for each sample defined in the `targets` file. This includes all the output and log files for each step. When these options are used, the output location will be updated by default and can be assigned to the same object.

```
## Run single Machine
args <- runCommandLine(args)
```

4.1.2 Parallelization on clusters

Alternatively, the computation can be greatly accelerated by processing many files in parallel using several compute nodes of a cluster, where a scheduling/queuing system is used for load balancing. For this the `clusterRun` function submits the computing requests to the scheduler using the run specifications defined by `runCommandLine`.

To avoid over-subscription of CPU cores on the compute nodes, the value from `yamlinput(args)['thread']` is passed on to the submission command, here `ncpus` in the `resources` list object. The number of independent parallel cluster processes is defined under the `Njobs` argument. The following example will run 18 processes in parallel using for each 4 CPU cores. If the resources available on a cluster allow running all 18 processes at the same time then the shown sample submission will utilize in total 72 CPU cores. Note, `clusterRun` can be used with most queueing systems as it is based on utilities from the `batchtools` package which supports the use of template files (`*.tpl`) for defining the run parameters of different schedulers. To run the following code, one needs to have both a conf file (see `.batchtools.conf.R` samples [here](#)) and a template file (see `*.tpl` samples [here](#)) for the queueing available on a system. The following example uses the sample conf and template files for the Slurm scheduler provided by this package.

```
library(batchtools)
resources <- list(walltime = 120, ntasks = 1, ncpus = 4, memory = 1024)
reg <- clusterRun(args, FUN = runCommandLine, more.args = list(args = args,
  make_bam = TRUE, dir = FALSE), conffile = ".batchtools.conf.R",
  template = "batchtools.slurm.tpl", Njobs = 18, runid = "01",
  resourceList = resources)
getStatus(reg = reg)
waitForJobs(reg = reg)
args <- output_update(args, dir = FALSE, replace = TRUE, extension = c(".sam",
  ".bam")) ## Updates the output(args) to the right location in the subfolders
output(args)
```

Check whether all BAM files have been created.

```
outpaths <- subsetWF(args, slot = "output", subset = 1, index = 1)
file.exists(outpaths)
```

4.2 Read and alignment stats

The following provides an overview of the number of reads in each sample and how many of them aligned to the reference.

```
read_statsDF <- alignStats(args = args)
write.table(read_statsDF, "results/alignStats.xls", row.names = FALSE,
            quote = FALSE, sep = "\t")
```

The following shows the alignment statistics for a sample file provided by the `systemPipeR` package.

```
read.table(system.file("extdata", "alignStats.xls", package = "systemPipeR"),
            header = TRUE)[1:4, ]
##   FileName Nreads2x Nalign Perc_Aligned Nalign_Primary
## 1      M1A   192918 177961    92.24697      177961
## 2      M1B   197484 159378    80.70426      159378
## 3      A1A   189870 176055    92.72397      176055
## 4      A1B   188854 147768    78.24457      147768
##   Perc_Aligned_Primary
## 1          92.24697
## 2          80.70426
## 3          92.72397
## 4          78.24457
```

4.3 Create symbolic links for viewing BAM files in IGV

The `symLink2bam` function creates symbolic links to view the BAM alignment files in a genome browser such as IGV. The corresponding URLs are written to a file with a path specified under `urlfile` in the `results` directory.

```
symLink2bam(sysargs = args, htmlDir = c("~/html/", "somedir/"),
            urlbase = "http://cluster.hpcc.ucr.edu/~tgirke/", urlfile = "./results/IGVurl.txt")
```

5 Read quantification

5.1 Read counting with `summarizeOverlaps` in parallel mode using multiple cores

Reads overlapping with annotation ranges of interest are counted for each sample using the `summarizeOverlaps` function (Lawrence et al. 2013). The read counting is preformed for exonic gene regions in a non-strand-specific manner while ignoring overlaps among different genes. Subsequently, the expression count values are normalized by *reads per kp per million mapped reads* (RPKM). The raw read count table (`countDFeByg.xls`) and the corresponding RPKM table (`rpkmDFeByg.xls`) are written to separate files in the directory of this project. Parallelization is achieved with the `BiocParallel` package, here using 8 CPU cores.

```
library("GenomicFeatures")
library(BiocParallel)
txdb <- makeTxDbFromGFF(file = "data/tair10.gff", format = "gff",
```



```

dataSource = "TAIR", organism = "Arabidopsis thaliana")
saveDb(txdb, file = "./data/tair10.sqlite")
txdb <- loadDb("./data/tair10.sqlite")
outpaths <- subsetWF(args, slot = "output", subset = 1, index = 1)
(aligned <- readGAlignments(outpaths[1])) # Demonstrates how to read bam file into R
eByg <- exonsBy(txdb, by = c("gene"))
bfl <- BamFileList(outpaths, yieldSize = 50000, index = character())
multicoreParam <- MulticoreParam(workers = 2)
register(multicoreParam)
registered()
counteByg <- bplapply(bfl, function(x) summarizeOverlaps(eByg,
  x, mode = "Union", ignore.strand = TRUE, inter.feature = FALSE,
  singleEnd = TRUE))
countDFeByg <- sapply(seq(along = counteByg), function(x) assays(counteByg[[x]])$counts)
rownames(countDFeByg) <- names(rowRanges(counteByg[[1]]))
colnames(countDFeByg) <- names(bfl)
rpkmDFeByg <- apply(countDFeByg, 2, function(x) returnRPKM(counts = x,
  ranges = eByg))
write.table(countDFeByg, "results/countDFeByg.xls", col.names = NA,
  quote = FALSE, sep = "\t")
write.table(rpkmDFeByg, "results/rpkmDFeByg.xls", col.names = NA,
  quote = FALSE, sep = "\t")

```

Sample of data slice of count table

```

read.delim("results/countDFeByg.xls", row.names = 1, check.names = FALSE)[1:4,
  1:5]

```

Sample of data slice of RPKM table

```

read.delim("results/rpkmDFeByg.xls", row.names = 1, check.names = FALSE)[1:4,
  1:4]

```

Note, for most statistical differential expression or abundance analysis methods, such as `edgeR` or `DESeq2`, the raw count values should be used as input. The usage of RPKM values should be restricted to specialty applications required by some users, e.g. manually comparing the expression levels among different genes or features.

5.2 Sample-wise correlation analysis

The following computes the sample-wise Spearman correlation coefficients from the `rlog` transformed expression values generated with the `DESeq2` package. After transformation to a distance matrix, hierarchical clustering is performed with the `hclust` function and the result is plotted as a dendrogram (also see file `sample_tree.pdf`).

```

library(DESeq2, quietly = TRUE)
library(ape, warn.conflicts = FALSE)
countDF <- as.matrix(read.table("./results/countDFeByg.xls"))
colData <- data.frame(row.names = targets.as.df(targets(args))$SampleName,
  condition = targets.as.df(targets(args))$Factor)
dds <- DESeqDataSetFromMatrix(countData = countDF, colData = colData,
  design = ~condition)

```

```
d <- cor(assay(rlog(dds)), method = "spearman")
hc <- hclust(dist(1 - d))
pdf("results/sample_tree.pdf")
plot.phylo(as.phylo(hc), type = "p", edge.col = "blue", edge.width = 2,
  show.node.label = TRUE, no.margin = TRUE)
dev.off()
```




Figure 2: Correlation dendrogram of samples

6 Analysis of DEGs

The analysis of differentially expressed genes (DEGs) is performed with the glm method of the `edgeR` package (Robinson, McCarthy, and Smyth 2010). The sample comparisons used by this analysis are defined in the header lines of the `targets.txt` file starting with `<CMP>`.

6.1 Run edgeR

```
library(edgeR)
countDF <- read.delim("results/countDFeByg.xls", row.names = 1,
  check.names = FALSE)
targets <- read.delim("targetsPE.txt", comment = "#")
cmp <- readComp(file = "targetsPE.txt", format = "matrix", delim = "-")
edgeDF <- run_edgeR(countDF = countDF, targets = targets, cmp = cmp[[1]],
  independent = FALSE, mdsplot = "")
```

Add gene descriptions

```
library("biomaRt")
m <- useMart("plants_mart", dataset = "athaliana_eg_gene", host = "plants.ensembl.org")
desc <- getBM(attributes = c("tair_locus", "description"), mart = m)
desc <- desc[!duplicated(desc[, 1]), ]
descv <- as.character(desc[, 2])
names(descv) <- as.character(desc[, 1])
edgeDF <- data.frame(edgeDF, Desc = descv[rownames(edgeDF)],
  check.names = FALSE)
write.table(edgeDF, "./results/edgeRglm_allcomp.xls", quote = FALSE,
  sep = "\t", col.names = NA)
```

6.2 Plot DEG results

Filter and plot DEG results for up and down regulated genes. The definition of *up* and *down* is given in the corresponding help file. To open it, type `?filterDEGs` in the R console.

```
edgeDF <- read.delim("results/edgeRglm_allcomp.xls", row.names = 1,
  check.names = FALSE)
pdf("results/DEGcounts.pdf")
DEG_list <- filterDEGs(degDF = edgeDF, filter = c(Fold = 2, FDR = 20))
dev.off()
write.table(DEG_list$Summary, "./results/DEGcounts.xls", quote = FALSE,
  sep = "\t", row.names = FALSE)
```




Figure 3: Up and down regulated DEGs with FDR of 1%

6.3 Venn diagrams of DEG sets

The `overLapper` function can compute Venn intersects for large numbers of sample sets (up to 20 or more) and plots 2-5 way Venn diagrams. A useful feature is the possibility to combine the counts from several Venn comparisons with the same number of sample sets in a single Venn diagram (here for 4 up and down DEG sets).

```
vennsetup <- overLapper(DEG_list$Up[6:9], type = "vennsets")
vennsetdown <- overLapper(DEG_list$Down[6:9], type = "vennsets")
pdf("results/vennplot.pdf")
vennPlot(list(vennsetup, vennsetdown), mymain = "", mysub = "",
  colmode = 2, ccol = c("blue", "red"))
dev.off()
```




Figure 4: Venn Diagram for 4 Up and Down DEG Sets

7 GO term enrichment analysis

7.1 Obtain gene-to-GO mappings

The following shows how to obtain gene-to-GO mappings from `biomaRt` (here for *A. thaliana*) and how to organize them for the downstream GO term enrichment analysis. Alternatively, the gene-to-GO mappings can be obtained for many organisms from Bioconductor's `*.db` genome annotation packages or GO annotation files provided by various genome databases. For each annotation this relatively slow preprocessing step needs to be performed only once. Subsequently, the preprocessed data can be loaded with the `load` function as shown in the next subsection.

```
library("biomaRt")
listMarts() # To choose BioMart database
listMarts(host = "plants.ensembl.org")
m <- useMart("plants_mart", host = "plants.ensembl.org")
```

```

listDatasets(m)
m <- useMart("plants_mart", dataset = "athaliana_eg_gene", host = "plants.ensembl.org")
listAttributes(m) # Choose data types you want to download
go <- getBM(attributes = c("go_id", "tair_locus", "namespace_1003"),
            mart = m)
go <- go[go[, 3] != "", ]
go[, 3] <- as.character(go[, 3])
go[go[, 3] == "molecular_function", 3] <- "F"
go[go[, 3] == "biological_process", 3] <- "P"
go[go[, 3] == "cellular_component", 3] <- "C"
go[1:4, ]
dir.create("./data/GO")
write.table(go, "data/GO/GOannotationsBiomart_mod.txt", quote = FALSE,
            row.names = FALSE, col.names = FALSE, sep = "\t")
catdb <- makeCATdb(myfile = "data/GO/GOannotationsBiomart_mod.txt",
                  lib = NULL, org = "", colno = c(1, 2, 3), idconv = NULL)
save(catdb, file = "data/GO/catdb.RData")

```

7.2 Batch GO term enrichment analysis

Apply the enrichment analysis to the DEG sets obtained the above differential expression analysis. Note, in the following example the FDR filter is set here to an unreasonably high value, simply because of the small size of the toy data set used in this vignette. Batch enrichment analysis of many gene sets is performed with the function. When `method=all`, it returns all GO terms passing the p-value cutoff specified under the `cutoff` arguments. When `method=slim`, it returns only the GO terms specified under the `myslimv` argument. The given example shows how a GO slim vector for a specific organism can be obtained from BioMart.

```

library("biomaRt")
load("data/GO/catdb.RData")
DEG_list <- filterDEGs(degDF = edgeDF, filter = c(Fold = 2, FDR = 50),
                      plot = FALSE)
up_down <- DEG_list$UporDown
names(up_down) <- paste(names(up_down), "_up_down", sep = "")
up <- DEG_list$Up
names(up) <- paste(names(up), "_up", sep = "")
down <- DEG_list$Down
names(down) <- paste(names(down), "_down", sep = "")
DEGlist <- c(up_down, up, down)
DEGlist <- DEGlist[sapply(DEGlist, length) > 0]
BatchResult <- GOcluster_Report(catdb = catdb, setlist = DEGlist,
                               method = "all", id_type = "gene", CLSZ = 2, cutoff = 0.9,
                               gocats = c("MF", "BP", "CC"), recordSpecGO = NULL)
library("biomaRt")
m <- useMart("plants_mart", dataset = "athaliana_eg_gene", host = "plants.ensembl.org")
goslimvec <- as.character(getBM(attributes = c("goslim_goa_accession"),
                                mart = m)[, 1])
BatchResultslim <- GOcluster_Report(catdb = catdb, setlist = DEGlist,
                                   method = "slim", id_type = "gene", myslimv = goslimvec, CLSZ = 10,
                                   cutoff = 0.01, gocats = c("MF", "BP", "CC"), recordSpecGO = NULL)

```

7.3 Plot batch GO term results

The `data.frame` generated by `GOCluster` can be plotted with the `goBarplot` function. Because of the variable size of the sample sets, it may not always be desirable to show the results from different DEG sets in the same bar plot. Plotting single sample sets is achieved by subsetting the input data frame as shown in the first line of the following example.

```
gos <- BatchResultslim[grep("M6-V6_up_down", BatchResultslim$CLID),
]
gos <- BatchResultslim
pdf("GOslimbarplotMF.pdf", height = 8, width = 10)
goBarplot(gos, gocat = "MF")
dev.off()
goBarplot(gos, gocat = "BP")
goBarplot(gos, gocat = "CC")
```




Figure 5: GO Slim Barplot for MF Ontology

8 Clustering and heat maps

The following example performs hierarchical clustering on the `rlog` transformed expression matrix subsetted by the DEGs identified in the above differential expression analysis. It uses a Pearson correlation-based distance measure and complete linkage for cluster joining.

```
library(pheatmap)
geneids <- unique(as.character(unlist(DEG_list[[1]])))
y <- assay(rlog(dds))[geneids, ]
pdf("heatmap1.pdf")
pheatmap(y, scale = "row", clustering_distance_rows = "correlation",
          clustering_distance_cols = "correlation")
dev.off()
```




Figure 6: Heat Map with Hierarchical Clustering Dendrograms of DEGs

9 Version Information

```
sessionInfo()
## R Under development (unstable) (2020-10-24 r79367)
## Platform: x86_64-pc-linux-gnu (64-bit)
## Running under: Ubuntu 20.04.1 LTS
##
```


RNA-Seq Workflow Template

```
## Matrix products: default
## BLAS: /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
## LAPACK: /home/dcassol/src/R-devel/lib/libRlapack.so
##
## locale:
## [1] LC_CTYPE=en_US.UTF-8 LC_NUMERIC=C
## [3] LC_TIME=en_US.UTF-8 LC_COLLATE=en_US.UTF-8
## [5] LC_MONETARY=en_US.UTF-8 LC_MESSAGES=en_US.UTF-8
## [7] LC_PAPER=en_US.UTF-8 LC_NAME=C
## [9] LC_ADDRESS=C LC_TELEPHONE=C
## [11] LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] stats4 parallel stats graphics grDevices
## [6] utils datasets methods base
##
## other attached packages:
## [1] batchtools_0.9.14 ape_5.4-1
## [3] ggplot2_3.3.2 systemPipeR_1.23.9
## [5] ShortRead_1.47.2 GenomicAlignments_1.25.3
## [7] SummarizedExperiment_1.19.9 Biobase_2.49.1
## [9] MatrixGenerics_1.1.8 matrixStats_0.57.0
## [11] BiocParallel_1.23.3 Rsamtools_2.5.3
## [13] Biostrings_2.57.2 XVector_0.29.3
## [15] GenomicRanges_1.41.6 GenomeInfoDb_1.25.11
## [17] IRanges_2.23.10 S4Vectors_0.27.14
## [19] BiocGenerics_0.35.4 BiocStyle_2.17.2
##
## loaded via a namespace (and not attached):
## [1] colorspace_1.4-1 rjson_0.2.20
## [3] hwriter_1.3.2 ellipsis_0.3.1
## [5] bit64_4.0.5 AnnotationDbi_1.51.3
## [7] xml2_1.3.2 codetools_0.2-16
## [9] splines_4.1.0 knitr_1.30
## [11] jsonlite_1.7.1 annotate_1.67.2
## [13] GO.db_3.12.0 dbplyr_1.4.4
## [15] png_0.1-7 pheatmap_1.0.12
## [17] graph_1.67.1 BiocManager_1.30.10
## [19] compiler_4.1.0 httr_1.4.2
## [21] GOstats_2.55.0 backports_1.1.10
## [23] assertthat_0.2.1 Matrix_1.2-18
## [25] limma_3.45.19 formatR_1.7
## [27] htmltools_0.5.0 prettyunits_1.1.1
## [29] tools_4.1.0 gtable_0.3.0
## [31] glue_1.4.2 GenomeInfoDbData_1.2.4
## [33] Category_2.55.0 dplyr_1.0.2
## [35] rsvg_2.1 rappdirs_0.3.1
## [37] V8_3.2.0 Rcpp_1.0.5
## [39] vctrs_0.3.4 nlme_3.1-149
## [41] rtracklayer_1.49.5 xfun_0.18
## [43] stringr_1.4.0 lifecycle_0.2.0
```

```
## [45] XML_3.99-0.5          edgeR_3.31.7
## [47] zlibbioc_1.35.0       scales_1.1.1
## [49] BSgenome_1.57.7       VariantAnnotation_1.35.4
## [51] hms_0.5.3             RBGL_1.65.0
## [53] RColorBrewer_1.1-2    yaml_2.2.1
## [55] curl_4.3              memoise_1.1.0
## [57] biomaRt_2.45.9        latticeExtra_0.6-29
## [59] stringi_1.4.6         RSQLite_2.2.1
## [61] genefilter_1.71.0     checkmate_2.0.0
## [63] GenomicFeatures_1.41.3 DDT_0.1
## [65] rlang_0.4.8           pkgconfig_2.0.3
## [67] bitops_1.0-6          evaluate_0.14
## [69] lattice_0.20-41       purrr_0.3.4
## [71] bit_4.0.4             tidyselect_1.1.0
## [73] GSEABase_1.51.3       AnnotationForge_1.31.3
## [75] magrittr_1.5          bookdown_0.21
## [77] R6_2.4.1              generics_0.0.2
## [79] base64url_1.4         DelayedArray_0.15.16
## [81] DBI_1.1.0             pillar_1.4.6
## [83] withr_2.3.0           survival_3.2-7
## [85] RCurl_1.98-1.2        tibble_3.0.4
## [87] crayon_1.3.4          BiocFileCache_1.13.1
## [89] rmarkdown_2.5         jpeg_0.1-8.1
## [91] progress_1.2.2        locfit_1.5-9.4
## [93] grid_4.1.0            data.table_1.13.2
## [95] blob_1.2.1            Rgraphviz_2.33.0
## [97] digest_0.6.27         xtable_1.8-4
## [99] brew_1.0-6            openssl_1.4.3
## [101] munsell_0.5.0         askpass_1.1
```

10 Funding

This project was supported by funds from the National Institutes of Health (NIH).

References

- H Backman, Tyler W, and Thomas Girke. 2016. "systemPipeR: NGS workflow and report generation environment." *BMC Bioinformatics* 17 (1): 388. <https://doi.org/10.1186/s12859-016-1241-0>.
- Howard, Brian E, Qiwen Hu, Ahmet Can Babaoglu, Manan Chandra, Monica Borghi, Xiaoping Tan, Luyan He, et al. 2013. "High-Throughput RNA Sequencing of Pseudomonas-Infected Arabidopsis Reveals Hidden Transcriptome Complexity and Novel Splice Variants." *PLoS One* 8 (10): e74183. <https://doi.org/10.1371/journal.pone.0074183>.
- Kim, Daehwan, Ben Langmead, and Steven L Salzberg. 2015. "HISAT: A Fast Spliced Aligner with Low Memory Requirements." *Nat. Methods* 12 (4): 357–60.

RNA-Seq Workflow Template

Lawrence, Michael, Wolfgang Huber, Hervé Pagès, Patrick Aboyoun, Marc Carlson, Robert Gentleman, Martin T Morgan, and Vincent J Carey. 2013. "Software for Computing and Annotating Genomic Ranges." *PLoS Comput. Biol.* 9 (8): e1003118. <https://doi.org/10.1371/journal.pcbi.1003118>.

Morgan, Martin, Hervé Pagès, Valerie Obenchain, and Nathaniel Hayden. 2019. *Rsamtools: Binary Alignment (Bam), Fasta, Variant Call (Bcf), and Tabix File Import*. <http://bioconductor.org/packages/Rsamtools>.

Robinson, M D, D J McCarthy, and G K Smyth. 2010. "EdgeR: A Bioconductor Package for Differential Expression Analysis of Digital Gene Expression Data." *Bioinformatics* 26 (1): 139–40. <https://doi.org/10.1093/bioinformatics/btp616>.